

Supervised Learning-Linear Regression

October 22, 2019

```
In [1]: from sklearn import datasets
iris = datasets.load_iris()
```

```
In [4]: print(iris.feature_names)
```

```
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
```

```
In [7]: print(iris.target)
         print(iris.target_names)
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
2 2]  
['setosa' 'versicolor' 'virginica']
```

```
In [8]: import pandas as pd
df = pd.DataFrame(iris.data)
```


```
print(df.head())
```

	0	1	2	3
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

```
In [9]: %matplotlib inline
        from matplotlib import pyplot as plt
        from sklearn.datasets.samples_generator import make_regression

        X, y = make_regression(n_samples=100, n_features=1, noise=5.4)
        plt.scatter(X,y)
```


Out[9]: <matplotlib.collections.PathCollection at 0x21831584208>



```
In [10]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
from sklearn.datasets import make_blobs

X, y = make_blobs(500, centers=3)
rgb = np.array(['r', 'g', 'b'])
plt.scatter(X[:, 0], X[:, 1], color=rgb[y])
```

Out[10]: <matplotlib.collections.PathCollection at 0x218325eb908>




```
In [13]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
from sklearn.datasets import make_circles

X, y = make_circles(n_samples=100, noise=0.09)

rgb=np.array(['r', 'g', 'b'])
plt.scatter(X[:, 0], X[:, 1], color=rgb[y])

Out[13]: <matplotlib.collections.PathCollection at 0x2183268f9b0>
```




```
In [15]: %matplotlib inline
import matplotlib.pyplot as plt
# represents the heights of a group of people in meters
heights = [[1.6], [1.65], [1.7], [1.73], [1.8]]

# represents the weights of a group of people in kgs
weights = [[60], [65], [72.3], [75], [80]]

plt.title('Weights plotted against heights')
plt.xlabel('Heights in meters')
plt.ylabel('Weights in kgs')

plt.plot(heights, weights, 'k.')

# axis range for x and y
plt.axis([1.5, 1.85, 50, 90])
plt.grid(True)
```



```
In [16]: from sklearn.linear_model import LinearRegression
         # Create and fit the model
         model = LinearRegression()
         model.fit(X=heights, y=weights)
```

```
Out[16]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                          normalize=False)
```

```
In [17]: # make prediction
         weight = model.predict([[1.75]])[0][0]
         print(round(weight,2))
```

76.04

```
In [18]: import matplotlib.pyplot as plt

         heights = [[1.6], [1.65], [1.7], [1.73], [1.8]]
         weights = [[60], [65], [72.3], [75], [80]]


         plt.title('Weights plotted against heights')
         plt.xlabel('Heights in meters')
         plt.ylabel('Weights in kilograms')
```

```
plt.plot(heights, weights, 'k.')

plt.axis([1.5, 1.85, 50, 90] )
plt.grid(True)

#plt the regression line
plt.plot(heights, model.predict(heights), color='r')
```

Out[18]: [<matplotlib.lines.Line2D at 0x2183273c080>]




```
In [22]: plt.title('Weights plotted against heights')
plt.xlabel('Heights in meters')
plt.ylabel('Weights in kilograms')

plt.plot(heights, weights, 'k.')

plt.axis([0, 1.85, -200, 200])
plt.grid(True)
#plot the regression line
extreme_heights = [[0], [1.8]]
plt.plot(extreme_heights, model.predict(extreme_heights), color='b')
```

Out[22]: [<matplotlib.lines.Line2D at 0x21832b232b0>]



```
In [24]: round(model.predict([[0]])[0][0],2)
```

```
Out[24]: -104.75
```

```
In [25]: print(round(model.intercept_[0],2))
```

```
-104.75
```

```
In [26]: print(round(model.coef_[0][0],2))
```

```
103.31
```

```
In [27]: import numpy as np
```

```
print('Residual sum of squares: %.2f' % np.sum((weights-model.predict(heights))**2))
```

```
Residual sum of squares: 5.34
```

```
In [29]: heights_test = [[1.58], [1.62], [1.69], [1.76], [1.82]]
weights_test = [[58], [63], [72], [73], [85]]
```

```
#Total Sum of squares (TSS)
```

```

weights_test_mean = np.mean(np.ravel(weights_test))
TSS = np.sum((np.ravel(weights_test)-weights_test_mean)**2)
print("TSS: %.2f" % TSS)
#Residual Sum of Squares (RSS)
RSS = np.sum((np.ravel(weights_test)-np.ravel(model.predict(heights_test))))**2)
print("RSS: %.2f" % RSS)

# R_squared
R_squared = 1-(RSS/TSS)
print("R-squared: %.2f" % R_squared)

```

TSS: 430.80

RSS: 24.62

R-squared: 0.94

```

In [30]: # using scikit-learn to calculates r-squared
print('R-squared: %.4f' % model.score(heights_test, weights_test))

```

R-squared: 0.9429

```

In [31]: import pickle

```

```

In [32]: # save the model to disk
filename = 'HeightsAndWeights_model.sav'
# write to the file using write and binary mode
pickle.dump(model, open(filename, 'wb'))

```

```

In [33]: # load the model from disk
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(heights_test, weights_test)

```

```

In [34]: print(result)

```

0.9428592885995253

```

In [35]: from sklearn.externals import joblib

```

```

In [36]: # save the model to disk
filename = 'HeightsAndWeights_model2.sav'
joblib.dump(model, filename)
# load the model from disk
loaded_model = joblib.load(filename)
result = loaded_model.score(heights_test, weights_test)
print(result)

```

0.9428592885995253


```
In [37]: # Clearing Rows with NaNs
```

```
import pandas as pd
df = pd.read_csv('NaNDataset.csv')
df.isnull().sum()
```

```
Out[37]: A    0
         B    2
         C    0
         dtype: int64
```

```
In [38]: # replaces all the NaNs in column B with the average of column B
```

```
df.B = df.B.fillna(df.B.mean())
print(df)
```

```
   A    B    C
0   1  2.0    3
1   4 11.0    6
2   7 11.0    9
3  10 11.0   12
4  13 14.0   15
5  16 17.0   18
```

```
In [39]: df = pd.read_csv('NaNDataset.csv')
```

```
df = df.dropna()
print(df)
```

```
   A    B    C
0   1  2.0    3
3  10 11.0   12
4  13 14.0   15
5  16 17.0   18
```

```
In [41]: df = df.reset_index(drop=True)
```

```
print(df)
```

```
   A    B    C
0   1  2.0    3
1  10 11.0   12
2  13 14.0   15
3  16 17.0   18
```

```
In [45]: import pandas as pd
```

```
df = pd.read_csv('DuplicateRows.csv')
print(df.duplicated(keep=False))
```

```
0    False
1     True
```

```

2      True
3     False
4     False
5      True
6      True
7     False
8     False
dtype: bool

```

```
In [47]: print(df[df.duplicated(keep=False)])
```

```

      A  B  C
1    4  5  6
2    4  5  6
5   10 11 12
6   10 11 12

```

```
In [51]: df.drop_duplicates(keep='first', inplace = True)
         print(df)
```

```

      A  B  C
0    1  2  3
1    4  5  6
3    7  8  9
4    7 18  9
5   10 11 12
7   13 14 15
8   16 17 18

```

```
In [53]: df.drop_duplicates(subset=['A', 'C'], keep='last', inplace = True)
         print(df)
```

```

      A  B  C
0    1  2  3
1    4  5  6
4    7 18  9
5   10 11 12
7   13 14 15
8   16 17 18

```

```
In [55]: import pandas as pd
         from sklearn import preprocessing
```

```
In [56]: df = pd.read_csv('NormalizeColumns.csv')
```

```
In [57]: x=df.values.astype(float)
        min_max_scaler = preprocessing.MinMaxScaler()
        x_scaled = min_max_scaler.fit_transform(x)
        df=pd.DataFrame(x_scaled, columns=df.columns)
        print(df)
```

	A	B	C
0	0.6	0.000000	0.0
1	0.2	0.200000	0.2
2	0.4	0.266667	0.4
3	0.0	0.600000	0.6
4	0.8	0.800000	0.8
5	1.0	1.000000	1.0

```
In [63]: import pandas as pd
        df = pd.read_csv("http://www.mosaic-web.org/go/datasets/galton.csv")
        print(df.head(20))
```

	family	father	mother	sex	height	nkids
0	1	78.5	67.0	M	73.2	4
1	1	78.5	67.0	F	69.2	4
2	1	78.5	67.0	F	69.0	4
3	1	78.5	67.0	F	69.0	4
4	2	75.5	66.5	M	73.5	4
5	2	75.5	66.5	M	72.5	4
6	2	75.5	66.5	F	65.5	4
7	2	75.5	66.5	F	65.5	4
8	3	75.0	64.0	M	71.0	2
9	3	75.0	64.0	F	68.0	2
10	4	75.0	64.0	M	70.5	5
11	4	75.0	64.0	M	68.5	5
12	4	75.0	64.0	F	67.0	5
13	4	75.0	64.0	F	64.5	5
14	4	75.0	64.0	F	63.0	5
15	5	75.0	58.5	M	72.0	6
16	5	75.0	58.5	M	69.0	6
17	5	75.0	58.5	M	68.0	6
18	5	75.0	58.5	F	66.5	6
19	5	75.0	58.5	F	62.5	6

```
In [60]: import numpy as np
        def outliers_iqr(data):
            q1, q3 = np.percentile(data, [25, 75])
            iqr = q3-q1
            lower_bound = q1-(iqr*1.5)
            upper_bound = q3+(iqr*1.5)
            return np.where((data>upper_bound)|(data<lower_bound))
```

```
In [61]: for i in outliers_iqr(df.height)[0]:
         print(df[i:i+1])
```

	family	father	mother	sex	height	nkids
288	72	70.0	65.0	M	79.0	7

```
In [64]: def outliers_z_score(data):
         threshold =3
         mean = np.mean(data)
         std = np.std(data)
         z_scores = [(y-mean)/std for y in data]
         return np.where(np.abs(z_scores)>threshold)
```

```
for i in outliers_z_score(df.height)[0]:
    print(df[i:i+1])
print()
```

	family	father	mother	sex	height	nkids
125	35	71.0	69.0	M	78.0	5
	family	father	mother	sex	height	nkids
288	72	70.0	65.0	M	79.0	7
	family	father	mother	sex	height	nkids
672	155	68.0	60.0	F	56.0	7

```
In [65]: # Supervised learning - Linear Regression
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

from sklearn.datasets import load_boston
dataset = load_boston()
```

```
In [66]: print(dataset)
```

```
{'data': array([[6.3200e-03, 1.8000e+01, 2.3100e+00, ..., 1.5300e+01, 3.9690e+02,
 4.9800e+00],
 [2.7310e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9690e+02,
 9.1400e+00],
 [2.7290e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9283e+02,
 4.0300e+00],
 ...,
 [6.0760e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
 5.6400e+00],
 [1.0959e-01, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9345e+02,
 6.4800e+00],
 [4.7410e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
```

```

7.8800e+00]]), 'target': array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5,
18.9, 21.7, 20.4, 18.2, 19.9, 23.1, 17.5, 20.2, 18.2, 13.6, 19.6,
15.2, 14.5, 15.6, 13.9, 16.6, 14.8, 18.4, 21. , 12.7, 14.5, 13.2,
13.1, 13.5, 18.9, 20. , 21. , 24.7, 30.8, 34.9, 26.6, 25.3, 24.7,
21.2, 19.3, 20. , 16.6, 14.4, 19.4, 19.7, 20.5, 25. , 23.4, 18.9,
35.4, 24.7, 31.6, 23.3, 19.6, 18.7, 16. , 22.2, 25. , 33. , 23.5,
19.4, 22. , 17.4, 20.9, 24.2, 21.7, 22.8, 23.4, 24.1, 21.4, 20. ,
20.8, 21.2, 20.3, 28. , 23.9, 24.8, 22.9, 23.9, 26.6, 22.5, 22.2,
23.6, 28.7, 22.6, 22. , 22.9, 25. , 20.6, 28.4, 21.4, 38.7, 43.8,
33.2, 27.5, 26.5, 18.6, 19.3, 20.1, 19.5, 19.5, 20.4, 19.8, 19.4,
21.7, 22.8, 18.8, 18.7, 18.5, 18.3, 21.2, 19.2, 20.4, 19.3, 22. ,
20.3, 20.5, 17.3, 18.8, 21.4, 15.7, 16.2, 18. , 14.3, 19.2, 19.6,
23. , 18.4, 15.6, 18.1, 17.4, 17.1, 13.3, 17.8, 14. , 14.4, 13.4,
15.6, 11.8, 13.8, 15.6, 14.6, 17.8, 15.4, 21.5, 19.6, 15.3, 19.4,
17. , 15.6, 13.1, 41.3, 24.3, 23.3, 27. , 50. , 50. , 50. , 22.7,
25. , 50. , 23.8, 23.8, 22.3, 17.4, 19.1, 23.1, 23.6, 22.6, 29.4,
23.2, 24.6, 29.9, 37.2, 39.8, 36.2, 37.9, 32.5, 26.4, 29.6, 50. ,
32. , 29.8, 34.9, 37. , 30.5, 36.4, 31.1, 29.1, 50. , 33.3, 30.3,
34.6, 34.9, 32.9, 24.1, 42.3, 48.5, 50. , 22.6, 24.4, 22.5, 24.4,
20. , 21.7, 19.3, 22.4, 28.1, 23.7, 25. , 23.3, 28.7, 21.5, 23. ,
26.7, 21.7, 27.5, 30.1, 44.8, 50. , 37.6, 31.6, 46.7, 31.5, 24.3,
31.7, 41.7, 48.3, 29. , 24. , 25.1, 31.5, 23.7, 23.3, 22. , 20.1,
22.2, 23.7, 17.6, 18.5, 24.3, 20.5, 24.5, 26.2, 24.4, 24.8, 29.6,
42.8, 21.9, 20.9, 44. , 50. , 36. , 30.1, 33.8, 43.1, 48.8, 31. ,
36.5, 22.8, 30.7, 50. , 43.5, 20.7, 21.1, 25.2, 24.4, 35.2, 32.4,
32. , 33.2, 33.1, 29.1, 35.1, 45.4, 35.4, 46. , 50. , 32.2, 22. ,
20.1, 23.2, 22.3, 24.8, 28.5, 37.3, 27.9, 23.9, 21.7, 28.6, 27.1,
20.3, 22.5, 29. , 24.8, 22. , 26.4, 33.1, 36.1, 28.4, 33.4, 28.2,
22.8, 20.3, 16.1, 22.1, 19.4, 21.6, 23.8, 16.2, 17.8, 19.8, 23.1,
21. , 23.8, 23.1, 20.4, 18.5, 25. , 24.6, 23. , 22.2, 19.3, 22.6,
19.8, 17.1, 19.4, 22.2, 20.7, 21.1, 19.5, 18.5, 20.6, 19. , 18.7,
32.7, 16.5, 23.9, 31.2, 17.5, 17.2, 23.1, 24.5, 26.6, 22.9, 24.1,
18.6, 30.1, 18.2, 20.6, 17.8, 21.7, 22.7, 22.6, 25. , 19.9, 20.8,
16.8, 21.9, 27.5, 21.9, 23.1, 50. , 50. , 50. , 50. , 50. , 13.8,
13.8, 15. , 13.9, 13.3, 13.1, 10.2, 10.4, 10.9, 11.3, 12.3, 8.8,
7.2, 10.5, 7.4, 10.2, 11.5, 15.1, 23.2, 9.7, 13.8, 12.7, 13.1,
12.5, 8.5, 5. , 6.3, 5.6, 7.2, 12.1, 8.3, 8.5, 5. , 11.9,
27.9, 17.2, 27.5, 15. , 17.2, 17.9, 16.3, 7. , 7.2, 7.5, 10.4,
8.8, 8.4, 16.7, 14.2, 20.8, 13.4, 11.7, 8.3, 10.2, 10.9, 11. ,
9.5, 14.5, 14.1, 16.1, 14.3, 11.7, 13.4, 9.6, 8.7, 8.4, 12.8,
10.5, 17.1, 18.4, 15.4, 10.8, 11.8, 14.9, 12.6, 14.1, 13. , 13.4,
15.2, 16.1, 17.8, 14.9, 14.1, 12.7, 13.5, 14.9, 20. , 16.4, 17.7,
19.5, 20.2, 21.4, 19.9, 19. , 19.1, 19.1, 20.1, 19.9, 19.6, 23.2,
29.8, 13.8, 13.3, 16.7, 12. , 14.6, 21.4, 23. , 23.7, 25. , 21.8,
20.6, 21.2, 19.1, 20.6, 15.2, 7. , 8.1, 13.6, 20.1, 21.8, 24.5,
23.1, 19.7, 18.3, 21.2, 17.5, 16.8, 22.4, 20.6, 23.9, 22. , 11.9]], 'feature_names': ar
'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7'), 'DESCR': ".. _boston_dataset:\n\nBoston I

```

```
In [68]: print(dataset.feature_names)
```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'  
 'B' 'LSTAT']
```

```
In [69]: print(dataset.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 506
```

```
:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is us
```

```
:Attribute Information (in order):
```

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

```
:Missing Attribute Values: None
```

```
:Creator: Harrison, D. and Rubinfeld, D.L.
```

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics

...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the AAAI Conference on Artificial Intelligence, pp. 466-471.

```
In [70]: print(dataset.target) # The prices of houses is the information we are seeking  
        # and it can be accessed via the target property
```

```
[24.  21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15.  18.9 21.7 20.4  
18.2 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8  
18.4 21.  12.7 14.5 13.2 13.1 13.5 18.9 20.  21.  24.7 30.8 34.9 26.6  
25.3 24.7 21.2 19.3 20.  16.6 14.4 19.4 19.7 20.5 25.  23.4 18.9 35.4  
24.7 31.6 23.3 19.6 18.7 16.  22.2 25.  33.  23.5 19.4 22.  17.4 20.9  
24.2 21.7 22.8 23.4 24.1 21.4 20.  20.8 21.2 20.3 28.  23.9 24.8 22.9  
23.9 26.6 22.5 22.2 23.6 28.7 22.6 22.  22.9 25.  20.6 28.4 21.4 38.7  
43.8 33.2 27.5 26.5 18.6 19.3 20.1 19.5 19.5 20.4 19.8 19.4 21.7 22.8  
18.8 18.7 18.5 18.3 21.2 19.2 20.4 19.3 22.  20.3 20.5 17.3 18.8 21.4  
15.7 16.2 18.  14.3 19.2 19.6 23.  18.4 15.6 18.1 17.4 17.1 13.3 17.8  
14.  14.4 13.4 15.6 11.8 13.8 15.6 14.6 17.8 15.4 21.5 19.6 15.3 19.4  
17.  15.6 13.1 41.3 24.3 23.3 27.  50.  50.  50.  22.7 25.  50.  23.8  
23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6 29.9 37.2 39.8 36.2  
37.9 32.5 26.4 29.6 50.  32.  29.8 34.9 37.  30.5 36.4 31.1 29.1 50.  
33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50.  22.6 24.4 22.5 24.4 20.  
21.7 19.3 22.4 28.1 23.7 25.  23.3 28.7 21.5 23.  26.7 21.7 27.5 30.1  
44.8 50.  37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29.  24.  25.1 31.5  
23.7 23.3 22.  20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8  
29.6 42.8 21.9 20.9 44.  50.  36.  30.1 33.8 43.1 48.8 31.  36.5 22.8  
30.7 50.  43.5 20.7 21.1 25.2 24.4 35.2 32.4 32.  33.2 33.1 29.1 35.1  
45.4 35.4 46.  50.  32.2 22.  20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9  
21.7 28.6 27.1 20.3 22.5 29.  24.8 22.  26.4 33.1 36.1 28.4 33.4 28.2  
22.8 20.3 16.1 22.1 19.4 21.6 23.8 16.2 17.8 19.8 23.1 21.  23.8 23.1  
20.4 18.5 25.  24.6 23.  22.2 19.3 22.6 19.8 17.1 19.4 22.2 20.7 21.1  
19.5 18.5 20.6 19.  18.7 32.7 16.5 23.9 31.2 17.5 17.2 23.1 24.5 26.6  
22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7 22.7 22.6 25.  19.9 20.8 16.8  
21.9 27.5 21.9 23.1 50.  50.  50.  50.  50.  13.8 13.8 15.  13.9 13.3  
13.1 10.2 10.4 10.9 11.3 12.3  8.8  7.2 10.5  7.4 10.2 11.5 15.1 23.2  
 9.7 13.8 12.7 13.1 12.5  8.5  5.  6.3  5.6  7.2 12.1  8.3  8.5  5.  
11.9 27.9 17.2 27.5 15.  17.2 17.9 16.3  7.  7.2  7.5 10.4  8.8  8.4  
16.7 14.2 20.8 13.4 11.7  8.3 10.2 10.9 11.  9.5 14.5 14.1 16.1 14.3  
11.7 13.4  9.6  8.7  8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
```

```

14.1 13. 13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20. 16.4 17.7
19.5 20.2 21.4 19.9 19. 19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3
16.7 12. 14.6 21.4 23. 23.7 25. 21.8 20.6 21.2 19.1 20.6 15.2 7.
8.1 13.6 20.1 21.8 24.5 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9
22. 11.9]

```

In [72]: *# load the dat into a Pandas DataFrame*

```

df = pd.DataFrame(dataset.data, columns=dataset.feature_names)
df.head()

```

```

Out [72]:      CRIM      ZN  INDUS  CHAS    NOX     RM   AGE     DIS  RAD    TAX  \
0  0.00632  18.0    2.31   0.0  0.538  6.575  65.2  4.0900  1.0  296.0
1  0.02731   0.0    7.07   0.0  0.469  6.421  78.9  4.9671  2.0  242.0
2  0.02729   0.0    7.07   0.0  0.469  7.185  61.1  4.9671  2.0  242.0
3  0.03237   0.0    2.18   0.0  0.458  6.998  45.8  6.0622  3.0  222.0
4  0.06905   0.0    2.18   0.0  0.458  7.147  54.2  6.0622  3.0  222.0

      PTRATIO      B  LSTAT
0      15.3  396.90   4.98
1      17.8  396.90   9.14
2      17.8  392.83   4.03
3      18.7  394.63   2.94
4      18.7  396.90   5.33

```

In [73]: *#add the prices of the houses to the DataFrame, add a new column to the DataFrame and*

```

df['MEDV'] = dataset.target
df.head()

```

```

Out [73]:      CRIM      ZN  INDUS  CHAS    NOX     RM   AGE     DIS  RAD    TAX  \
0  0.00632  18.0    2.31   0.0  0.538  6.575  65.2  4.0900  1.0  296.0
1  0.02731   0.0    7.07   0.0  0.469  6.421  78.9  4.9671  2.0  242.0
2  0.02729   0.0    7.07   0.0  0.469  7.185  61.1  4.9671  2.0  242.0
3  0.03237   0.0    2.18   0.0  0.458  6.998  45.8  6.0622  3.0  222.0
4  0.06905   0.0    2.18   0.0  0.458  7.147  54.2  6.0622  3.0  222.0

      PTRATIO      B  LSTAT  MEDV
0      15.3  396.90   4.98  24.0
1      17.8  396.90   9.14  21.6
2      17.8  392.83   4.03  34.7
3      18.7  394.63   2.94  33.4
4      18.7  396.90   5.33  36.2

```

In [74]: *#Data cleansing*

```

df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):

```



```

CRIM      506 non-null float64
ZN        506 non-null float64
INDUS     506 non-null float64
CHAS      506 non-null float64
NOX       506 non-null float64
RM        506 non-null float64
AGE       506 non-null float64
DIS       506 non-null float64
RAD       506 non-null float64
TAX       506 non-null float64
PTRATIO   506 non-null float64
B         506 non-null float64
LSTAT     506 non-null float64
MEDV      506 non-null float64
dtypes: float64(14)
memory usage: 55.5 KB

```

```
In [75]: print(df.isnull().sum())
```

```

CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
MEDV      0
dtype: int64

```

```

In [76]: # Feature selection
         # choose those features that directly influence the result(that is prices of houses)
         # corr() function computes the pairwise correlation of columns
corr = df.corr()
print(corr)

```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	\
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	

RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339
MEDV	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955

	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
CRIM	-0.379670	0.625505	0.582764	0.289946	-0.385064	0.455621	-0.388305
ZN	0.664408	-0.311948	-0.314563	-0.391679	0.175520	-0.412995	0.360445
INDUS	-0.708027	0.595129	0.720760	0.383248	-0.356977	0.603800	-0.483725
CHAS	-0.099176	-0.007368	-0.035587	-0.121515	0.048788	-0.053929	0.175260
NOX	-0.769230	0.611441	0.668023	0.188933	-0.380051	0.590879	-0.427321
RM	0.205246	-0.209847	-0.292048	-0.355501	0.128069	-0.613808	0.695360
AGE	-0.747881	0.456022	0.506456	0.261515	-0.273534	0.602339	-0.376955
DIS	1.000000	-0.494588	-0.534432	-0.232471	0.291512	-0.496996	0.249929
RAD	-0.494588	1.000000	0.910228	0.464741	-0.444413	0.488676	-0.381626
TAX	-0.534432	0.910228	1.000000	0.460853	-0.441808	0.543993	-0.468536
PTRATIO	-0.232471	0.464741	0.460853	1.000000	-0.177383	0.374044	-0.507787
B	0.291512	-0.444413	-0.441808	-0.177383	1.000000	-0.366087	0.333461
LSTAT	-0.496996	0.488676	0.543993	0.374044	-0.366087	1.000000	-0.737663
MEDV	0.249929	-0.381626	-0.468536	-0.507787	0.333461	-0.737663	1.000000

```
In [77]: print(df.corr().abs().nlargest(3, 'MEDV').index)#top 3 correlation index
```


```
Index(['MEDV', 'LSTAT', 'RM'], dtype='object')
```

```
In [79]: print(df.corr().abs().nlargest(3, 'MEDV').values[:,13])#top3 correlation values
```

```
[1.          0.73766273 0.69535995]
```


```
In [80]: #Since RM and LSTAT have high correlation values, we will use these two features to t
#Multiple Regression
# 2 or more independent variables are used to predict the value of a dependent variab
# Plot a scatter plot showing the relationship between the LSTAT feature and the MEDV
plt.scatter(df['LSTAT'], df['MEDV'], marker='o')
plt.xlabel('LSTAT')
plt.ylabel('MEDV')
```

```
Out[80]: Text(0, 0.5, 'MEDV')
```




```
In [81]: # RM and MEDV label
plt.scatter(df['RM'], df['MEDV'], marker='o')
plt.xlabel('RM')
plt.ylabel('MEDV')
```

```
Out[81]: Text(0, 0.5, 'MEDV')
```



```
In [82]: from mpl_toolkits.mplot3d import Axes3D
```

```
In [83]: fig = plt.figure(figsize=(18,15))
ax=fig.add_subplot(111, projection='3d')
ax.scatter(df['LSTAT'], df['RM'], df['MEDV'], c='b')
ax.set_xlabel("LSTAT")
ax.set_ylabel("RM")
ax.set_zlabel("MEDV")
plt.show()
```



```
In [88]: # Training the model
x = pd.DataFrame(np.c_[df['LSTAT'], df['RM']], columns=['LSTAT', 'RM'])
Y = df['MEDV']
```

```
In [89]: from sklearn.model_selection import train_test_split
x_train, x_test, Y_train, Y_test = train_test_split(x, Y, test_size=0.3, random_state=
```

```

In [90]: print(x_train.shape)
          print(Y_train.shape)

(354, 2)
(354,)

In [91]: # x training set 354rows, 2col    Y 354rows 1col
          print(x_test.shape)
          print(Y_test.shape)

(152, 2)
(152,)

In [93]: from sklearn.linear_model import LinearRegression
          model = LinearRegression()
          model.fit(x_train, Y_train)

Out[93]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                          normalize=False)

In [94]: price_pred = model.predict(x_test)

In [95]: print('R-Squared: %.4f' % model.score(x_test, Y_test))

R-Squared: 0.6162


In [96]: from sklearn.metrics import mean_squared_error
          mse = mean_squared_error(Y_test, price_pred)
          print(mse)

36.49422110915324

In [97]: plt.scatter(Y_test, price_pred)
          plt.xlabel("Actual Prices")
          plt.ylabel('Predicted prices')
          plt.title("Actual prices vs Predicted prices")

Out[97]: Text(0.5, 1.0, 'Actual prices vs Predicted prices')

```



```
In [98]: print(model.intercept_)
         print(model.coef_)
```

```
0.38437936780346504
[-0.65957972  4.83197581]
```

```
In [100]: print(model.predict([[30, 5]]))
```

```
[4.75686695]
```

```
In [103]: import matplotlib.pyplot as plt
         import pandas as pd
         import numpy as np
         from mpl_toolkits.mplot3d import Axes3D
         from sklearn.datasets import load_boston
         dataset = load_boston()
         df = pd.DataFrame(dataset.data, columns=dataset.feature_names)
         df['MEDV'] = dataset.target
```

```
In [108]: x = pd.DataFrame(np.c_[df['LSTAT'], df['RM']], columns=['LSTAT', 'RM'])
         Y = df['MEDV']
         fig = plt.figure(figsize=(18,15))
```

```

ax = fig.add_subplot(111, projection='3d')
ax.scatter(x['LSTAT'], x['RM'], Y, c='b')
ax.set_xlabel("LSTAT")
ax.set_ylabel("RM")
ax.set_zlabel("MEDV")

x_surf = np.arange(0,40,1)
y_surf = np.arange(0,10,1)
x_surf, y_surf = np.meshgrid(x_surf, y_surf)


from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x, Y)

z = lambda x,y: (model.intercept_+model.coef_[0] * x + model.coef_[1]*y)

ax.plot_surface(x_surf, y_surf, z(x_surf,y_surf),
               rstride=1,
               cstride=1,
               color='None',
               alpha=0.4)

plt.show()

```



```


In [109]: # Polynomial Regression
df = pd.read_csv('polynomial.csv')
plt.scatter(df.x,df.y)

```


```

Out[109]: <matplotlib.collections.PathCollection at 0x2183323ee80>

```



```
In [111]: x=df.x[0:6, np.newaxis]# convert to 2D array
          y=df.y[0:6, np.newaxis]
          model.fit(x,y)
          # perform prediction
          y_pred = model.predict(x)
          # plot the training point
          plt.scatter(x, y, s=10, color='b')
          # plot the straight line
          plt.plot(x, y_pred, color='r')
          plt.show()
          # calculate R-Squared
          print('R-Squared for training set: %.4f' % model.score(x,y))
```

R-Squared for training set: 0.8658

```
In [129]: from sklearn.preprocessing import PolynomialFeatures
          degree=5
          polynomial_features = PolynomialFeatures(degree=degree)
          x_poly=polynomial_features.fit_transform(x)
          print(x_poly)


[[1.00000e+00 1.50000e+00 2.25000e+00 3.37500e+00 5.06250e+00 7.59375e+00]
 [1.00000e+00 2.00000e+00 4.00000e+00 8.00000e+00 1.60000e+01 3.20000e+01]
 [1.00000e+00 3.00000e+00 9.00000e+00 2.70000e+01 8.10000e+01 2.43000e+02]
 [1.00000e+00 4.00000e+00 1.60000e+01 6.40000e+01 2.56000e+02 1.02400e+03]
 [1.00000e+00 5.00000e+00 2.50000e+01 1.25000e+02 6.25000e+02 3.12500e+03]
 [1.00000e+00 6.00000e+00 3.60000e+01 2.16000e+02 1.29600e+03 7.77600e+03]]
```

```
In [130]: print(polynomial_features.get_feature_names('x'))

['1', 'x', 'x^2', 'x^3', 'x^4', 'x^5']
```

```
In [131]: model = LinearRegression()
          model.fit(x_poly, y)
          y_poly_pred = model.predict(x_poly)
          # plot the points
```

```
plt.scatter(x,y, s=10)
# plot the regression line
plt.plot(x, y_poly_pred)
plt.show()
print(model.intercept_)
print(model.coef_)
```



```
[14.02380952]
[[ 0.         -27.35119048  20.692791    -6.71693122   0.99371693
  -0.05489418]]
```

```
In [132]: print('R-squared for training set: % .4f' % model.score(x_poly, y))
```

```
R-squared for training set:  1.0000
```

```
In [133]: # Using Polynomial Multiple Regression on the Boston Dataset
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.datasets import load_boston
```

```

In [135]: dataset = load_boston()
          df = pd.DataFrame(dataset.data, columns = dataset.feature_names)
          df['MEDV'] = dataset.target
          x=pd.DataFrame(np.c_[df['LSTAT'], df['RM']], columns=['LSTAT', 'RM'])
          Y=df['MEDV']
          from sklearn.model_selection import train_test_split
          x_train, x_test, Y_train, Y_test = train_test_split(x, Y, test_size=0.3, random_state=42)

In [136]: degree = 2
          polynomial_features = PolynomialFeatures(degree = degree)
          x_train_poly = polynomial_features.fit_transform(x_train)

In [138]: print(polynomial_features.get_feature_names(['x', 'y']))

['1', 'x', 'y', 'x^2', 'x y', 'y^2']

In [140]: model = LinearRegression()
          model.fit(x_train_poly, Y_train)

Out[140]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                          normalize=False)

In [141]: x_test_poly = polynomial_features.fit_transform(x_test)
          print('R-squared: %.4f' % model.score(x_test_poly, Y_test))

R-squared: 0.7340

In [142]: print(model.intercept_)
          print(model.coef_)

26.93343052383913
[ 0.00000000e+00  1.47424550e+00 -6.70204730e+00  7.93570743e-04
 -3.66578385e-01  1.17188007e+00]

In [143]: import matplotlib.pyplot as plt
          import pandas as pd
          import numpy as np
          from mpl_toolkits.mplot3d import Axes3D
          from sklearn.preprocessing import PolynomialFeatures
          from sklearn.linear_model import LinearRegression
          from sklearn.datasets import load_boston

In [145]: dataset = load_boston()
          df['MEDV']=dataset.target
          x=pd.DataFrame(np.c_[df['LSTAT'], df['RM']], columns = ['LSTAT', 'RM'])
          Y=df['MEDV']

```

```

In [150]: fig=plt.figure(figsize=(18,15))
ax=fig.add_subplot(111, projection='3d')
ax.scatter(x['LSTAT'],x['RM'],Y,c='b')
ax.set_xlabel("LSTAT")
ax.set_ylabel("RM")
ax.set_zlabel("MEDV")


x_surf = np.arange(0,40,1)
y_surf = np.arange(0,10,1)
x_surf, y_surf = np.meshgrid(x_surf, y_surf)

degree=2
polynomial_features = PolynomialFeatures(degree=degree)
x_poly = polynomial_features.fit_transform(x)
print(polynomial_features.get_feature_names(['x','y']))
model = LinearRegression()
model.fit(x_poly, Y)
z=lambda x, y: (model.intercept_+
                (model.coef_[1]* x)+
                (model.coef_[2]* y)+
                (model.coef_[3]* x**2)+
                (model.coef_[4]* x*y)+
                (model.coef_[5]* y**2)
                )
ax.plot_surface(x_surf, y_surf, z(x_surf, y_surf),
               rstride=1,
               cstride=1,
               color='None',
               alpha=0.4)

plt.show()

['1', 'x', 'y', 'x^2', 'x y', 'y^2']

```



```

In [ ]:

```